

CSE 40625 — Machine Learning

Assignment 4 (10 points)

Due Date: March. 30, 2017 (Sakai Drop Box)

Single-Layer Neural Network

Overview

The purpose of this assignment is for you to implement a single-layer neural network, a supervised machine learning inspired by the way the brain works. Neural networks are systems of interconnected “neurons” or units that can compute values from inputs by feeding information through the network. The model makes predictions by chaining together layers of neural units. A single-layer neural network has a single layer of hidden neurons. Materials for the assignment, including the dataset, expected output, and template code can be found [here](#) on GitHub.

You may use Python libraries for handling data preprocessing and visualization, including but not limited to NumPy, SciPy, pandas, and Matplotlib, but *you may NOT use any Python libraries that employ machine learning models, including but not limited to scikit-learn, StatsModels, TensorFlow, or Orange*. Your solution to the assignment should be individually submitted.

Dataset

You will use the “digits” dataset on handwritten digit classification with all 10 classes (labeled from 0 to 9) for this assignment. The data is provided in comma-separated (CSV) file format. For all rows, the last column designates the class (y) and the remaining columns designate features (X). The first row consists of the feature and class names.

Procedure

Use the above digits dataset as input to a single-layer neural network model with 100 hidden units. Initialize the weights at each layer with random samples drawn from a uniform distribution over $[-0.5, 0.5]$. Initialize the activations as input x . Train the network with batches of 100 instances at a time. For each batch of inputs x , feed the activations of the batch forward through the network. Add a fixed bias term of positive one to the activations. At the output layer, apply the softmax tangent function to compute the output predictions. Using multinomial logistic loss (also known as cross-entropy loss), compute the minimize gradient of the error function. Update the weights at each layer using the gradient with a learning rate of 0.01. Repeat for at 500 iterations. Predict each target value based on the unit in the output layer with the highest activation output.

Output

Your code should output the accuracy of the prediction results every 50 iterations while fitting the model, a blank line, and the confusion matrix of the final output results.

Example output:

```
0 0.845
50 0.931
100 0.932
150 0.931
200 0.931
250 0.931
300 0.932
350 0.932
400 0.933
450 0.933
```

Predicted	0	1	2	3	4	5	6	7	8	9
Actual										
0	521	20	2	0	2	7	0	0	1	1
1	0	532	4	0	0	0	0	0	19	16
2	0	15	518	9	0	0	1	1	11	2
3	0	3	1	538	0	8	0	3	13	6
4	1	18	0	0	520	0	8	2	6	13
5	0	12	1	7	2	517	0	0	9	10
6	0	15	1	0	5	0	525	0	12	0
7	0	8	0	1	4	4	0	532	11	6
8	0	23	0	1	1	1	0	0	524	4
9	0	13	0	8	8	2	0	6	7	518

Submission

Please submit a Python executable (singlelayer_neuralnet.py) file of your code to your Sakai Drop Box.

Should you run into any problems, please feel free to email or meet with the instructor.